

# SHUBHASHIS SINHA

Senior Staff AI Engineer — Enterprise AI Platform · RAG & Agents · MCP / MCP Hub · Bedrock AgentCore · Microservices/APIs

✉ shubhashis.sinha@gmail.com | 📞 +91-7358604077 | 📍 Bangalore, India | 🔗 [linkedin.com/in/shubhashis-sinha](https://linkedin.com/in/shubhashis-sinha)

---

August 5, 2026

Dear GE Vernova Hiring Team,

I am writing regarding the **Senior Staff AI Engineer** role in Digital Engineering (Bengaluru). As a current **Staff AI Architect at GE HealthCare's AI Center of Excellence**, I bring both a deeply relevant skill set and genuine familiarity with the GE engineering culture, standards, and ways of working. With **12+ years** leading enterprise full-stack and AI platform development at scale — shipping **microservices and REST APIs, RAG and multimodal LLM features, and agents**, and owning technical decisions end-to-end — my background maps closely to this role.

My GE HealthCare work is a strong proxy for what you're building. I led enterprise AI platform development across multiple initiatives — designing **agentic workflows for autonomous tasks** on LangGraph with **MCP** tool interfaces, building a GraphRAG multi-agent platform, and shipping a Contract Digitalization system where agents process **1,000+ documents/month**. Two capabilities I'd highlight as differentiators: I standardized agent tooling through an **MCP Hub** — a centralized registry of reusable MCP servers/tools shared across agents and teams — and I've deployed and scaled agentic workloads on **Amazon Bedrock AgentCore** (runtime, memory, gateway, identity) for secure, managed enterprise deployment. I built **enterprise RAG architectures** with semantic-search optimization and conversation memory over 10,000+ documents at sub-second P95, and delivered **reusable AI services** adopted across teams. Having done this inside GE, I understand how to deliver measurable impact within GE's scale, governance, and quality expectations — and I'd value continuing that contribution within the GE family.

I own the full engineering stack this role calls for: resilient **REST APIs and microservices** (FastAPI), full-stack delivery integrating backend AI services with front-end interfaces, and strong discipline around **reliability, performance, and cloud cost**. I set **standards for AI-assisted development** (daily use of agentic coding tools), establish **testing for prompts and model performance** (RAGAS, LLM-as-judge; a validator agent reached ~95% correctness), and implement **observability for AI features** with tracing and token/latency/cost monitoring. My cloud is AWS with Docker and Kubernetes. On front-end, I build and integrate React/TypeScript interfaces for AI products, while my deepest strength is backend and AI-platform engineering.

I also lead the way this role expects — mentoring engineers, guiding external contractors and vendors, running code reviews and integrations, and driving Agile delivery with autonomy and measurable impact. My 12+ years place me squarely within your requirements.

Continuing my journey within the GE family — now helping GE Vernova build enterprise AI platforms that advance the energy transition — is a genuinely exciting next step. I would welcome the opportunity to discuss how my agentic, RAG, and AI-platform experience can accelerate your initiatives.

Sincerely,

**Shubhashis Sinha**